Business Proposal: Stop Customers From Leaving

A Platform for Growth-Stage B2B SaaS

November 2025

Executive Summary

The Problem Nobody Talks About

Growth-stage B2B SaaS companies ($5-50M ARR) are bleeding customers. 35% annual churn is killing their economics, but they can't afford the $50K-100K/year enterprise platforms like Gainsight. And even if they could, those tools take 3 months to implement and need dedicated data teams.

Here's the truth: Nobody wakes up wanting "churn prediction." They want to **keep customers who are about to leave**. Prediction is just the mechanism. The value is in stopping the bleeding.

Right now, they're stuck. Most just have endless meetings where they try stuff without any systematic approach. Maybe a VP of Customer Success builds a spreadsheet. Maybe they set up some basic alerts in their CRM. But nobody has real predictive capability, and when they finally see a customer at risk, it's usually too late.

What We're Building

A self-service churn prediction platform that actually works. Here's what it does:

Data pipeline that sets itself up. Airbyte connects to Stripe, Salesforce, Segment, Zendesk - whatever they're already using. 600+ pre-built connectors mean we don't have to write custom code for every customer. They authenticate with OAuth, data flows automatically, and it's done quickly. Ideally, we'd be able to wrap a self-service version of Airbyte in a white-label wrapper.

Real ML models, not just "health scores." Ensemble models (scikit-learn + XGBoost + LightGBM) that predict churn 90 days out with 85%+ accuracy. Way better than "customer hasn't logged in for 30 days = red flag" rules.

Actual recommendations, not just alerts. This is the differentiator. We don't just say "Customer X is at risk." We say: "Customer X hasn't used Feature Y. Here's the email to send. Here's the call script. Reach out today, follow up Thursday, schedule a QBR next week." CS teams know precisely what to do.

Pricing that aligns with results. Base subscription ($200-2K/month) plus performance bonuses when we save customers. They only pay us extra when we can prove we prevent churn. De-risks the whole purchase.

Market Size

The customer success software market is $2.2B in 2025, growing 22% annually. But here's what matters: there are several thousand growth-stage B2B SaaS companies globally in the $5-50M ARR range. That's our target.

Most of them (I'd estimate 70%) could use a self-service product. The other 30% either have enough budget for Gainsight or need white-glove professional services. We're after the 70%.

By Year 3, we want 200 customers at roughly $1,500/month average. That's $3.6M ARR. Conservative but achievable with a bootstrap approach.

Why This Works

Fast implementation. For customers with standard tools (Stripe, Salesforce, HubSpot, Segment, Zendesk), you're live in 3-5 days. Connect data sources via OAuth in 30 minutes, historical data loads overnight, model trains in 24 hours, dashboard goes live. No professional services needed. Competitors like Gainsight take 8-12 weeks and require implementation consultants. If your data stack is mostly custom or you're missing a data warehouse, that's a different conversation - but if you're running a modern SaaS stack, this is plug and play.

Cheap enough to justify. $800-2,000/month vs $2,000-8,000/month for enterprise platforms. Even at their smallest, a growth-stage company can swing this.

ROI is obvious. Customer has $15M ARR, 35% churn? That's $5.25M lost revenue per year. We reduce churn by just 10% (35% → 31.5%)? They save $525K in revenue, $367K in gross profit. We cost them maybe $76K/year total. That's 380% ROI. Easy sell.

Self-service really means self-service. Airbyte + PostgreSQL + scikit-learn + Metabase. All open source or commodity infrastructure. We're not rebuilding data pipelines from scratch for each customer. Set up once, replicate infinitely.

The Numbers (Bootstrap Path)

	Year 1	Year 2	Year 3
Customers	30	100	200
ARR	$450K	$1.5M	$3.6M
Burn	-$817K	-$1.1M	-$1.1M
Breakeven	—	—	Q2 Year 4

We'll need about $2.5-3M in funding total to get to profitability. Raise $1M in Year 1 (angels), another $1.5-2M in Year 2 (small seed). Then we're generating enough cash to grow without more capital.

Unit economics work: 5.9:1 LTV:CAC, 6-month payback. That's healthy.

Market Analysis

How Big Is This Really?

Customer success software overall: $2.2B market in 2025, heading to $6B by 2030. Growing at 22% annually because every SaaS company now has a VP of Customer Success and a budget for tools.

But who actually needs what we're building?

There are roughly 30,000 SaaS companies worldwide, with about 17,000 in the United States. Around 8,500 of these have crossed $1M in annual revenue. Only about 13% of SaaS startups ever reach $10M ARR.

Here's the honest truth: nobody publishes reliable data on how many companies are in the $5-50M ARR sweet spot. We know it's somewhere between "a few thousand" and "several thousand" globally. The exact number matters less than this: it's a large enough market to build a real business, and it's dramatically underserved by current solutions.

What we know about these companies:

- They have 500-5,000 active customers paying $5K-20K/year on average
- Churn is killing them - 35% annually is typical
- CS team is 3-15 people, stretched thin
- They already spend $6K-25K/year on customer success platforms like Vitally, ChurnZero, Totango, or Custify. Budget exists. The question is whether they add us alongside their existing CS platform (we do prediction + remediation, they do workflow management), replace a tool that isn't working, or keep trying to build something themselves in spreadsheets (which usually fails after 6 months).
- They're too big for spreadsheets, too small for Gainsight

Our target: 200 customers in 3 years

Whether the addressable market is 3,000 companies or 8,000 companies, we're going after 200 of them. That's a tiny slice. If we can't find 200 growth-stage B2B SaaS companies who need help reducing churn, we have bigger problems than market sizing.

What's their situation?

They have 500-2,000 active customers paying $8K-15K/year on average. Churn rate is about 35% annually (worse than larger companies). CS team is 2-8 people, stretched thin. Maybe 1 CSM per 150-200 accounts. They're firefighting, not being strategic.

They already spend $12K-40K/year on customer success tools. Budget exists. The question is whether they spend it on us vs competitors vs building something themselves (which usually fails after 6 months).

Our Target: 200 Customers in 3 Years

Year 1: 30 customers (0.6% penetration of North America) Year 2: 100 customers (2.0% penetration) Year 3: 200 customers (4.0% penetration)

That's conservative. We're not assuming we capture 20% of the market. We're assuming slow, steady growth with a bootstrap budget and sales-assisted PLG motion.

If we nail product-market fit and raise more capital, could we grow faster? Sure. But let's prove the model first.

Customer Value Prop

What Problem Are We Actually Solving?

Primary issue: They don't have sophisticated churn prediction and have no idea what to do when someone's at risk.

Why the problem exists:

Data is a mess. Customer data lives in 8-12 different tools. Nobody has built a data warehouse because that requires engineering resources they don't have. Even if they had clean data, they don't have ML expertise to build models.

No actionability. Even companies that see early warning signs (usage dropping, support tickets spiking, NPS declining) don't know what specific action to take. Send a generic "check-in" email? That doesn't work. The CS team is winging it.

Resource constraints. 1 CSM managing 150-200 accounts can't be proactive. They're responding to whoever emails them or whose renewal is coming up. There's no system.

What They're Doing Today (The Alternatives)

Most common (40% of market): Spreadsheet analysis VP of CS exports data from Stripe, Salesforce, Segment once a month. Looks for patterns. Identifies "red flag" accounts based on gut feel. Maybe creates a list of 20 customers to reach out to. It's manual, time-consuming, and they're catching problems 30 days before churn (too late).

Second most common (30%): BI dashboards Someone set up Metabase or Looker. They can see trailing indicators (who churned last month, what cohorts have high churn). But no predictive capability. Just reporting on what already happened.

Some teams (20%): Basic health scores in their CRM Salesforce or HubSpot has simple rules: "If login frequency < 2x/week AND support tickets > 3 AND usage trending down = red." Better than nothing, but crude. No ML, no playbooks.

The rest (10%): Nothing systematic CS team is completely reactive. Customer cancels? "Oh no, let's try to save them." Obviously too late.

Our Solution in Plain English

Week 1-2: Connect your data Customer signs up, clicks through OAuth screens to connect Stripe (billing), Salesforce (CRM), Segment (product usage), Zendesk (support). No IT involvement. Data starts flowing. We provision a PostgreSQL database with their own schema.

Load 6-12 months of historical data.

Week 3-4: Train models, start predicting Our ML pipeline extracts features (days since last login, average session time, feature usage patterns, support ticket sentiment, payment failures). Trains ensemble models. Within a week, every customer has a churn risk score updated daily. Dashboard goes live.

Week 4+: Get recommendations Customer X has 75% churn risk. Dashboard shows: "Customer X hasn't used Feature Y in 45 days. Competitors who use Feature Y have 2x lower churn. Action: Send this email template introducing Feature Y, offer 15-min training call, follow up in 3 days if no response."

CSM clicks button, email template populates with customer details, sends. Three days later, automated reminder. A week later, schedule a QBR. System tracks whether they acted on the recommendation and whether it worked.

Quarterly: Strategic insights Not just individual customer predictions. We analyze cohorts: "Customers who don't complete onboarding in 7 days churn at 2x the rate. Recommend updating onboarding flow." Or: "60% of churned customers cited missing Feature X. Consider adding this to roadmap."

That's it. The platform is doing the heavy lifting. CS team just has to execute the playbooks.

Show Me the ROI

Let's use a realistic example:

Baseline scenario: - Company: $15M ARR - Customers: 1,000 active accounts - Average contract value: $15,000 - Churn rate: 35% annually - Gross margin: 70%

What churn costs them: - Lost customers: 350 per year - Lost revenue: $5.25M - Lost gross profit: $3.68M

With 10% churn reduction (35% → 31.5%): - Save 35 additional customers - Retain $525K more revenue - Keep $367K more gross profit

What we cost: - Platform: $24K/year ($2,000/month Tier 3) - Performance bonuses: ~$52K (35 customers × $15K × 10%) - Total: $76K

Customer's net operational benefit: $291K (380% ROI in year one)

But that's just the P&L impact. Here's what really matters:

The Enterprise Value Impact

Venture-backed SaaS companies are valued on ARR multiples. Current market: private SaaS companies trade at **5-6x ARR** (public comps are 6-7x, private companies discount ~1x for illiquidity). Growth-stage companies with strong retention can command the higher end of that range.

Year 1 impact: - Retained ARR: $525K - Enterprise value created: $525K × 5.5x = **$2.9M**

But it compounds. Churn reduction isn't a one-time save - those customers keep paying year after year. And you're reducing churn on a growing base.

3-Year cumulative impact:

Year	Baseline ARR	With Platform	ARR Difference	Cumulative Retained
1	$15.0M	$15.0M	—	—
2	$12.6M	$13.4M	+$800K	$800K
3	$10.6M	$12.0M	+$1.4M	$2.2M

Assumes 35% churn baseline vs 25% with platform (10 point reduction), no new customer acquisition for simplicity

By Year 3: - **ARR difference: $1.4M higher** - **Enterprise value difference: $1.4M × 5.5x = $7.7M** - **Cumulative ARR retained: $2.2M** - **Cumulative enterprise value created: $12M+**

Why This Matters for Your Equity

The person buying this platform isn't just a VP with a retention target. They're an equity holder in a venture-backed company that will eventually exit - via acquisition, IPO, or (unfortunately) wind-down.

At a growth-stage company ($15-50M ARR), a VP/Head of Customer Success typically holds **0.5-2% equity**. Let's say 1%.

Without the platform: - Company valued at $15M × 5.5x = $82.5M - Their equity worth: $825K

With the platform (Year 3): - Company valued at $16.4M × 5.5x = $90.2M
- Their equity worth: $902K - **Personal gain: $77K**

And that's conservative - just the churn impact, ignoring that better retention often commands a

higher multiple because investors pay up for predictable revenue.

For the founder/CEO with 20% equity? That's a **$1.5M difference** in their personal outcome.

The real pitch: We're not selling a $24K/year subscription. We're selling a path to $2-3M higher enterprise value per year. The subscription is a rounding error.

For Venture-Backed Companies: The Clock is Ticking

Here's the thing about VC-backed companies: there's an exit timeline. The fund has a 10-year life. Your Series A was 3 years ago. You've got maybe 4-5 years to either:

1. Get acquired
2. Go public
3. Raise another round at a higher valuation
4. Run out of runway

In all of these scenarios, **ARR and retention rate are the two numbers that matter most.**

Net Revenue Retention (NRR) above 100% is table stakes for a premium valuation. Every point of churn reduction moves you toward that benchmark. Buyers and investors will pay more for a company with 85% gross retention than one with 75% - not just because the ARR is higher, but because the *multiple* is higher.

A company with strong retention might trade at 6-7x. One with leaky retention? 3-4x. Same ARR, half the valuation.

This is the pitch: Your churn problem isn't just costing you revenue. It's costing you millions in enterprise value. It's costing your team their equity upside. And the clock is running.

We fix it for $76K/year.

Competitive Landscape

Who Else Is Doing This?

Enterprise platforms (not really our competition):

Gainsight - The 800-pound gorilla. $26K-100K+/year. Takes 8-12 weeks to implement with professional services. Comprehensive features, great for enterprise. Total overkill for a $10M

ARR company. They can't afford it and wouldn't use half the features.

ChurnZero - Mid-market focused. $18K-60K/year. Still expensive, still slow (5-7 weeks setup). Better than Gainsight for our segment but still too much.

Mid-market platforms (partial overlap):

Vitally - This is probably our closest competitor. $14K-22K/year, fast setup (2-3 weeks), great UX. Weakness: Their ML is basic (rules-based health scores), and they don't have our remediation playbooks. They alert you to risk but don't tell you what to do.

Catalyst/Totango - $10K-25K/year. Decent feature set but hits scalability walls quickly. Basic churn prediction (not real ML). Good enough for some customers, not sophisticated enough for others.

What About Building In-House?

Lots of companies try this. VP of CS convinces CEO to dedicate 1-2 engineers for 3-6 months. They build a data pipeline, some dashboards, maybe a simple model.

6 months later, it's broken. Engineers get pulled to other projects. Data sources changed their APIs. The model needs retraining but nobody knows how. Project dies.

We've seen this pattern over and over. That's why we exist.

Where We Win

Time to value: We're live in 2-4 weeks. Gainsight takes 8-12 weeks. Time matters when churn is bleeding you right now.

Price: $200-2K/month vs $2K-8K/month for competitors. We're 50-75% cheaper. Growth-stage companies feel that difference.

Self-service: You sign up, connect data sources, we're running. No professional services, no implementation team, no "success manager" you have to schedule calls with. It just works.

Remediation focus: Competitors tell you who's at risk. We tell you what to do about it. That's the whole game.

Outcome-based pricing: We only get performance bonuses when we save customers. Risk-sharing. Nobody else does this.

Open source foundation: Airbyte, scikit-learn, PostgreSQL, Metabase. Standard, proven tools.

No proprietary lock-in. Customers can export their data anytime. We're not holding anything hostage.

What We Don't Do: The "Sleeping Dog" Problem

Here's something most churn tools get wrong: **low usage ≠ high churn risk.**

Some customers are happily paying for something they rarely use. It's a small line item nobody audits. They're keeping it "just in case." Enterprise contracts where procurement approved it and it just auto-renews.

If you flag these customers as "at risk" and send a check-in email, you remind them they're paying for something they don't use. They cancel. **You caused the churn you were trying to prevent.**

We call this the "sleeping dog" problem. Don't wake them up.

How we handle it:

Customer Profile	Usage	Trend	Our Recommendation
Active & healthy	High	Stable/up	No action needed
Active & declining	High	Down	**Intervene now**
Dormant & stable	Low	Flat 6+ months	**Leave alone**
Dormant & declining	Low	Down	**Intervene now**
New & struggling	Low	N/A	**Onboarding help**

Our models are trained on **actual churn events**, not just engagement proxies. We weight behavioral *changes* more heavily than absolute levels. A customer who logged in twice last year and twice this year isn't at risk - they're stable. A customer who logged in 50 times last month and 10 times this month? That's a problem.

Why this matters for performance bonuses:

Without this distinction, we could game the system: - Flag every low-usage customer as "at risk"
- Do minimal outreach
- Claim credit when they don't cancel (they were never going to) - Collect bonuses we didn't earn

That's bad ethics and bad business. Customers would figure it out. Instead, we only count saves where the customer showed **declining behavior** before intervention. Stable dormant accounts aren't eligible for performance bonuses. This keeps our incentives honest.

Where We're Weaker

Feature breadth: Gainsight has onboarding automation, in-app messaging, QBR scheduling, success planning, the works. We're laser-focused on churn prediction + remediation. That's intentional - we're specialists - but some buyers want an all-in-one suite.

Brand: They've been around 10+ years. We're new. Enterprise buyers feel safer with known brands. Growth-stage buyers care less about this.

Relationship selling: Big platforms assign you a CSM, do quarterly business reviews, have account executives. We're mostly self-service with support via email/chat. Some buyers want the high-touch relationship.

Enterprise features: SSO, SOC2, custom SLAs, on-premise deployment. We'll add these eventually as we move upmarket, but they're not in MVP.

Can these be overcome? Yes. - Feature breadth: Partner with existing CS tools, integrate via APIs - Brand: 100+ happy customers and strong case studies - Relationship selling: Use consultant partners for customers who want high-touch - Enterprise features: Add incrementally in Year 2-3

What Happens When Competitors React?

Gainsight launches "Gainsight Lite" at $1,000/month - Probably 12-18 months away (big companies move slow) - By then we have 100+ customers and strong positioning - Our self-service UX will still be better (they'll just scale down features)

Vitally adds better ML models - More likely, maybe 6-12 months - That's fine - we still have remediation playbooks as differentiation - And our pricing model (performance bonuses) is unique

Generic ML platforms (DataRobot, H2O.ai) add "churn prediction templates" - Always a risk - But they require technical expertise to use - We're turnkey for non-technical CS teams - Different buyer

The real moat develops over time: - We get better models as we see more customer data - We build a library of A/B tested playbooks (what actually works) - Customers connect 8-12

integrations to us (switching cost) - First-mover advantage in growth-stage segment

We’re not worried about competition. We’re worried about execution.

Product & Technical Architecture

MVP Feature Set

Phase 1: Core pipeline + basic prediction (Months 1-6)

Data integration layer:

- Pre-built connectors for the big 5: Stripe, Salesforce/HubSpot, Segment/Mixpanel, Zendesk/Intercom, product database (PostgreSQL/MySQL)
- OAuth authentication (no API keys to manage)
- PostgreSQL warehouse with per-tenant schemas (tenant1.customers, tenant2.customers)
- Historical data load: 6-12 months back

ML prediction:

- Logistic regression baseline (proven 86% accuracy)
- Features: days since last login, session duration, feature usage patterns, support ticket volume + sentiment, payment failures
- Predicts 90-day churn probability
- Daily batch scoring runs
- Flag customers with ≥70% risk

Basic dashboard:

- Embedded Metabase showing high-risk customers sorted by score
- Trend charts (risk over time)
- Export to CSV

- Email digest: daily alert of newly at-risk customers
- Slack integration

That's the MVP. Get this working, prove it's accurate, sign 5 alpha customers.

Phase 2: Remediation playbooks (Months 3-8)

This is where we differentiate:

Rule-based recommendations:

- If Feature X unused for 30+ days → "Send training email"
- If support tickets trending up → "Schedule check-in call"
- If payment method failing → "Update billing immediately"
- If usage declining 50%+ → "Escalate to VP CS"

Email templates:

- Pre-written with personalization tokens
- "Hi {{FirstName}}, I noticed you haven't tried {{FeatureName}} yet…"
- Customer just clicks "Send"

Call discussion guides:

- 3-5 talking points per at-risk customer
- "Ask about: their experience with X, whether they're achieving Y goal, roadblocks to adoption"
- Objection handling: common reasons for churn + responses

Intervention timeline:

- "Day 1: Send email A"
- "Day 3: If no response, call"
- "Day 7: Follow-up email B"

- “Day 14: Schedule QBR”

Customer exclusion:

- Let customers remove accounts from churn list (maybe they have inside info we don’t)
- Only provide recommendations for remaining accounts
- Track exclusion reasons for our learning

Performance tracking:

- Which customers were flagged high-risk?
- Which ones survived 4+ months?
- Calculate performance bonuses owed
- Quarterly reconciliation report

Phase 3: Advanced ML + strategic insights (Months 6-12)

Better models:

- Add Random Forest + XGBoost to ensemble
- Weighted average of predictions
- Target 90%+ accuracy
- A/B test model versions

Quarterly strategic reports:

- Cohort analysis: “Customers from Channel X churn at 2x rate vs Channel Y”
- Product gaps: “Top 3 missing features cited by churned customers”
- Training needs: “Help doc Y has 3x lower engagement than doc Z”
- Support patterns: “Ticket type A predicts 60% churn probability”

These insights help fix systemic issues, not just save individual customers.

More integrations:

- Expand to 15 connectors (Pendo, Amplitude, Helpscout, Drift, etc.)
- CSV upload for custom data
- REST API for programmable access

The Technical Stack

Data pipeline:

- Airbyte (open source): 600+ connectors, handles schema drift, CDC support
- dbt: SQL transformations, feature engineering
- PostgreSQL or Snowflake/BigQuery: Data warehouse (depends on customer size)
- Per-tenant schema isolation for security

ML layer:

- Python with scikit-learn, XGBoost, LightGBM
- MLflow: Experiment tracking, model versioning
- Batch scoring: Daily runs, results stored in warehouse
- REST API for on-demand predictions

Dashboard:

- Metabase (open source): Embedded, white-labeled at higher tiers
- Customer-facing UI: React frontend
- Backend: Python (FastAPI)

Orchestration:

- Apache Airflow or Prefect: Schedule jobs, coordinate pipelines

Hosting:

- AWS or GCP, Kubernetes for auto-scaling

- Separate S3 buckets per tenant for data sovereignty

Multi-Tenant Architecture

We're not giving every customer their own infrastructure. That doesn't scale.

Shared infrastructure, logical separation:

- All customers on same Kubernetes cluster
- Shared Airbyte instance, shared ML compute
- Separate schemas in the data warehouse (tenant1*., tenant2.*)
- Row-level security where needed
- Separate S3 buckets for sensitive data

Why this works:

- Cost-efficient (shared compute)
- Fast onboarding (provision new schema in seconds)
- Easy upgrades (apply to all tenants simultaneously)
- Can move to database-per-tenant for enterprise customers if needed

---

Cost Per Customer

Infrastructure costs:

- Data warehouse: $50-150/month (depends on data volume)
- Airbyte Cloud (if not self-hosted): $100-200/month
- Compute (ML training/scoring): $20-50/month
- Metabase Cloud (if not self-hosted): $0-20/month
- **Total: $170-420/month per customer**

At $800-2,000/month subscription, that's 75-80% gross margin. Good SaaS economics.

Open Source Advantage

Why we're using open source everywhere we can:

Lower costs: Airbyte Community Edition is free. PostgreSQL is free. scikit-learn is free. We're paying for hosting, not software licenses.

No vendor lock-in: Customer doesn't like us? They can export their data warehouse and build on it themselves. We're not holding anything hostage.

Battle-tested: These tools run at thousands of companies. They're proven, stable, well-documented.

Community support: Something breaks? Google it, find 10 Stack Overflow answers. vs proprietary tools where you're at the vendor's mercy.

Extensibility: Customer needs a custom connector? We can build it. vs closed-source tools where you have to beg the vendor for features.

The trade-off: We have to maintain infrastructure ourselves. But that's fine - we're a tech company. We should be good at ops.

Business Model & Pricing

The Tiers

Tier 1: Starter - $200/month

For small CS teams just getting started:

- Up to 500 customer records
- 3 data integrations (Stripe + CRM + one more)
- Daily churn scoring

- Email alerts
- 50 recommendations/month ($2 per additional)
- Email support (48hr response)
- Performance bonus eligible

Target customer: Early growth-stage, 1-2 person CS team, limited budget

Tier 2: Professional - $800/month

This is where most customers will land:

- Up to 2,000 customer records
- 8 data integrations
- Advanced ML (ensemble models)
- Slack integration
- Quarterly strategic reports
- 200 recommendations/month ($1.50 per additional)
- Priority support (email + chat, 24hr response)
- Performance bonus eligible

Target customer: Core growth-stage companies, 3-6 person CS team

Tier 3: Enterprise - $2,000/month

For larger customers or those who want white-glove:

- Up to 10,000 records
- Unlimited integrations
- White-labeled dashboards (remove our branding)
- 1 custom integration per quarter
- Monthly strategic reports

- Unlimited recommendations
- Dedicated success manager
- Phone support (8hr response)
- Quarterly business review
- Performance bonus eligible

Target customer: Upper end of growth-stage, 7+ person CS team, approaching mid-market

Tier 4: Custom

Call us. Includes:

- Database-per-tenant isolation (instead of shared)
- SSO/SAML
- Custom SLAs
- On-premise deployment
- API access
- Whatever else you need

Performance Bonus Structure

Here's how it works:

Trigger: We flag a customer with ≥70% churn risk score.

Success: That customer is still active 4 months later.

Payment: We get 10% of that customer's ACV, one time.

Frequency limit: Max once per customer per 12 months (prevents double-counting).

Customer control: They can remove accounts from our prediction list (maybe they have inside info). Those exclusions aren't eligible for bonuses.

Example:

- Company has 1,000 customers, $12K average ACV
- We flag 30 as high-risk in Q1
- 20 of those survive 4+ months
- Performance bonus: 20 × $12K × 10% = $24K
- Customer saved: 20 × $12K × 70% margin = $168K gross profit
- Our total fees: $24K subscription + $24K bonus = $48K
- Their net benefit: $120K (250% ROI)

Why 4 months?

- Too short (1-2 months) and we're not really proving we saved them
- Too long (6+ months) and we're waiting forever to get paid
- 4 months feels right - enough time to prove impact, not so long that customers forget

How do we avoid disputes? Clear contract language. Pre-agreed measurement methodology. Quarterly reconciliation report showing:

- Customer ID
- Flagged date
- Risk score
- ACV at time of flagging
- Status at 4-month mark
- Bonus due (if applicable)

Customer has audit rights. They can verify the list. If there's a disagreement, we discuss it. Relationship matters more than squeezing every dollar.

Add-On Revenue

Extra integrations: $100/month per connector beyond tier limits (only for standard Airbyte connectors)

Custom connector development: $2,000 upfront + $200/month maintenance (for specialized tools we don't have connectors for)

Professional services:

- Data source consultation: $200/hour (helping them clean up their data)
- Custom model tuning: $5,000 project (training models on their specific use case)
- Strategic CS consulting: $300/hour (Tier 1-2 only; included in Tier 3)

API access:

- Read API (export predictions): $500/month
- Write API (push custom data): $1,000/month

We're not building a big professional services business. These are just options for customers who want extra help. Maybe 5-10% of revenue.

Unit Economics

Customer Acquisition Cost: $5,000

Breakdown:

- Marketing spend: $1,500 (content, ads, demand gen)
- Sales time: $2,000 (demos, trial support, closing)
- Trial costs: $500 (infrastructure during trial period)
- Onboarding: $1,000 (activation support)

That's blended average. Organic/inbound customers might be $3K. Outbound might be $7K. We're estimating $5K overall.

Lifetime Value:

Average customer at $1,200/month (mix of Tier 2-3):

- Monthly revenue: $1,200
- Gross margin: 75%
- Gross profit/month: $900
- Expected lifetime: 33 months (90% annual retention = 10% monthly churn)
- **LTV: $900 × 33 = $29,700**

LTV:CAC = 5.9:1

That's healthy. SaaS benchmark is 3:1 minimum. We're at 6:1.

Payback period: 5.6 months

$5,000 CAC / $900 monthly gross profit = 5.6 months

Benchmark is <12 months good, <18 months acceptable. We're at 5.6 months. Excellent.

Why does this work?

- Low infrastructure costs (open source, shared infrastructure)
- Self-service reduces sales overhead
- High retention creates long LTV
- Performance bonuses are mostly margin (minimal cost to us)

Go-To-Market Strategy

Sales Model: Sales-Assisted PLG

We're not pure product-led growth. Deal size ($2.4K-24K/year) justifies some sales touch. Complex product (data + ML) benefits from human help. VP Customer Success is making a bet-your-job decision - they want to talk to someone.

But we're also not doing traditional enterprise sales. No custom demos, no 3-month evaluation cycles, no RFP processes. Customers can complete the entire journey online. Pricing is fixed

(no negotiation for Tiers 1-3).

The hybrid approach:

Most leads (70%) start self-service:

- Read our content, sign up for free trial
- Automated onboarding emails
- In-app guidance for connecting data sources
- Dashboard activates after first data sync
- “Book time with expert” button if they get stuck

But we touch every trial:

- Day 3: “How’s it going? Any blockers?” (15-min call)
- Day 7: “Let me show you advanced features” (20-min demo)
- Day 14: “Ready to move forward?” (30-min close call)

Top 20% of leads get white-glove treatment:

- Alpha/beta customers
- Reference account candidates
- Bigger deal size potential ($24K+)
- We’ll do extra onboarding support, custom walkthroughs, QBRs

---

Acquisition Channels

Year 1: Content + Inbound

We’re not spending big on paid ads yet. Build organic engine first.

SEO content targeting CS leaders:

- “How to Calculate Customer Churn Rate”

- “10 Early Warning Signs Your Customer Is About to Leave”
- “Build vs Buy: Customer Success Software Decision Framework”
- Target: 10,000 organic visits/month by Month 12

Founder-led content:

- LinkedIn posts (3x/week) on churn reduction tactics
- Guest posts on CS blogs (ChurnKey, Custify, Gainsight’s blog)
- Speaking at conferences (Customer Success Summit, Pulse)

Free tools as lead magnets:

- Churn rate calculator
- CAC payback calculator
- Customer health score template spreadsheet
- Gated downloads generate emails

Budget: $3K/month for content writers, SEO tools

Target: 50 qualified leads/month by Month 12

Year 2: Add Paid Acquisition

Once organic is working, layer on paid:

Google Ads:

- “Churn prediction software,” “customer success platform,” “reduce SaaS churn”
- Budget: $5K/month
- Target: 20 leads/month at $250 CPA

LinkedIn Ads:

- Target job titles: VP Customer Success, Director CS, CCO
- Target companies: $5-50M ARR B2B SaaS

- Budget: $3K/month
- Target: 10 leads/month at $300 CPA

Retargeting:

- Website visitors who viewed pricing
- Blog readers (3+ articles)
- Budget: $2K/month

Total paid: $10K/month for 30 additional qualified leads/month

Year 2-3: Partnerships

CS consultants:

- Find 10-20 independent consultants who advise growth-stage SaaS companies
- Offer 15-20% recurring commission on referrals
- They handle complex implementations, we provide the platform
- Target: 20% of customers via partners by Year 3

Integration partnerships:

- Co-marketing with Slack, Intercom, Zendesk, Segment
- Get listed in their app marketplaces
- Target: 10% of customers via integrations by Year 3

Community sponsorships:

- Sponsor Customer Success Slack groups
- Free accounts for CS influencers
- Co-create case studies

The Sales Process

Typical journey (45 days):

Week 1: Awareness

- Prospect finds us via blog post or Google search
- Downloads “Churn Reduction Playbook” (gated)
- Enters email nurture sequence

Weeks 2-4: Consideration

- Gets 5 educational emails over 3 weeks
- Reads comparison articles (us vs Gainsight vs ChurnZero vs build in-house)
- Talks to their VP CS peers (word-of-mouth)

Weeks 5-6: Trial (14 days)

- Signs up for free trial
- Day 1-2: Connects data sources (self-service)
- Day 3: Sales call #1 - “Any blockers? Let me help”
- Days 3-7: Historical data loads, models train
- Day 7: First predictions appear
- Day 7: Sales call #2 - “Here’s what we’re seeing”
- Days 8-14: Customer tests recommendations, shows to team
- Day 14: Sales call #3 - “Ready to go forward?”

Week 7-8: Close

- Tier 1-2: Online checkout (credit card)
- Tier 3: Annual invoice (NET 30 terms)
- Month-to-month for <$24K/year
- Annual contract (10% discount) for $24K+

Weeks 8-11: Onboarding

- Tier 1-2: Email onboarding, help docs
- Tier 3: Kick-off call with dedicated success manager
- Goal: First recommendation acted on within 30 days

Conversion metrics we're targeting:

- Trial signup → Paid: 20% (industry benchmark is 15-25%)
- Sales cycle: 30-45 days
- Average deal size: $1,200/month ($14,400/year)
- Close rate of qualified leads: 30%

Marketing Budget

Year 1: $73K

- Content: $36K
- Tools: $12K
- Events: $10K
- Website: $15K

Year 2: $241K

- Content: $48K
- Paid ads: $120K
- Tools: $18K
- Events: $25K
- Partnerships: $30K

Year 3: $324K

- Content: $60K
- Paid ads: $180K
- Tools: $24K
- Events: $40K
- Partnerships: $20K

That gives us blended CAC of $5K (marketing + sales costs together).

Reality Check: Customer Business Failures

The Blindspot Nobody Wants to Talk About

The industry talks about 35% annual churn like it's all fixable. It's not.

Here's what the data shows: 92% of SaaS startups fail within 3 years. 45% fail by year 5. 65% fail by year 10. Even after 10 years in business, only 13% of SaaS startups reach $10M ARR.

Our target customers are growth-stage companies ($5-50M ARR). They've survived early-stage, so they're not in the "92% fail in 3 years" bucket. But they're still vulnerable. 67% of SaaS companies experience their first major scaling crisis between $5M-$15M ARR.

Realistic Churn Model

Industry baseline (without us): 35% annual churn broken down as:

- **Customer business failures: 8%** - Companies going out of business, getting acquired, pivoting away
- **Controllable churn: 27%** - Product dissatisfaction, switching to competitors, budget cuts, poor onboarding

With our platform: 18% annual churn broken down as:

- **Customer business failures: 8%** - Still happens, outside our control

- **Controllable churn: 10%** - We reduce this by 63%

That's the value prop: **We can't save customers whose businesses fail, but we're experts at saving everyone else.**

Updated Unit Economics

With 18% annual churn (82% retention):

- Expected customer lifetime: 5.6 years
- LTV: $46,000
- CAC: $5,000
- LTV:CAC ratio: 9.2:1 (excellent)
- CAC payback: 7.2 months

This is healthy SaaS economics. The 9:1 LTV:CAC ratio gives us room for mistakes and shows the business model works even with realistic churn.

Financial Projections (Revised)

The Moderate Growth Path

With 18% annual churn, here's what it takes to reach ~200 customers by Year 3:

Year	New Customers	Churned	Ending Customers	Ending ARR
1	40	0	40	$442K
2	100	7	133	$1,468K
3	100	23	209	$2,307K

What this means:

- We acquire 240 total customers over 3 years to end with 209

- We lose 31 customers to churn (mostly in Y2-Y3 as the base grows)
- Total CAC spend: $1.2M ($5K × 240 customers)
- Blended MRR: $920/month

Year 1

Revenue:

- Average customers through year: 20 (ramping from 0 to 40)
- Subscription revenue: $221K
- Performance bonuses: $33K
- Total revenue: $254K

Expenses:

- COGS (25%): $64K
- Gross profit: $190K (75% margin)
- Sales & Marketing: $223K
- Engineering: $540K (3 FTEs @ $180K)
- Product: $160K (1 FTE)
- CS/Support: $100K (1 FTE)
- G&A: $32K
- Operating expenses: $1,055K
- **EBITDA: -$865K**

Burn rate: $72K/month

Year 2

Revenue:

- Average customers through year: 87 (from 40 to 133)
- Subscription revenue: $954K
- Performance bonuses: $143K
- Total revenue: $1,098K

Expenses:

- COGS (25%): $274K
- Gross profit: $824K (75% margin)
- Sales & Marketing: $591K
- Engineering: $900K (5 FTEs)
- Product: $320K (2 FTEs)
- CS/Support: $200K (2 FTEs)
- G&A: $138K
- Operating expenses: $2,149K
- **EBITDA: -$1,325K**

Burn rate: $110K/month

Customers lost to churn: 7 (18% of 40)

Year 3

Revenue:

- Average customers through year: 171 (from 133 to 209)
- Subscription revenue: $1,888K
- Performance bonuses: $283K

- Total revenue: $2,171K

Expenses:

- COGS (25%): $543K
- Gross profit: $1,628K (75% margin)
- Sales & Marketing: $824K
- Engineering: $1,440K (8 FTEs)
- Product: $480K (3 FTEs)
- CS/Support: $400K (4 FTEs)
- G&A: $314K
- Operating expenses: $3,458K
- **EBITDA: -$1,830K**

Burn rate: $152K/month

Customers lost to churn: 23 (18% of 133)

3-Year Summary

Metric	Year 1	Year 2	Year 3
Customers	40	133	209
ARR	$442K	$1,468K	$2,307K
Revenue	$254K	$1,098K	$2,171K
Gross Profit	$190K	$824K	$1,628K
EBITDA	-$865K	-$1,325K	-$1,830K
Cumulative Burn	$865K	$2,190K	$4,020K

Key metrics:

- Total customers acquired: 240
- Total CAC spend: $1.2M
- Cumulative 3-year burn: $4.0M
- Net Revenue Retention: ~100% (82% retention + some upsells)

Path to Profitability

With 18% annual churn, breakeven happens in **Q4 Year 5:**

Year 4 projections:

- Add 150 new customers, lose 38 to churn (18% of 209)
- Ending: 321 customers
- ARR: $3.5M
- EBITDA: -$2.3M (still burning)

Year 5 projections:

- Add 150 new customers, lose 58 to churn (18% of 321)
- Ending: 413 customers
- ARR: $4.6M
- Revenue: $5.3M
- Gross profit: $4.0M
- Operating expenses: $4.0M
- **EBITDA: breakeven**

From Year 6 onwards, we're profitable and can grow without raising more capital.

Total funding needed: ~$7M

- Seed (Year 1): $1.5M

- Series A (Year 2-3): $3.0M
- Series B (Year 4-5): $2.5M

This is realistic for a B2B SaaS company with strong unit economics (9:1 LTV:CAC). The 5-year path to profitability is longer than we'd like, but it reflects the reality of customer business failures we can't control.

Alternative Business Models We Considered

Pure Performance-Based (No Base Fee)

Structure: No subscription. We only charge 20-25% of retained revenue from saved customers.

Pros: Completely de-risks purchase for customer. Perfect incentive alignment. Premium positioning.

Cons: Zero revenue for first 4-6 months while proving value. Complex attribution. Hard to forecast. Needs $5M+ funding for 18-month runway.

Verdict: Too risky for bootstrap. Maybe viable later with more capital.

Freemium

Structure: Free tier for 100 customers, basic features. Paid tiers start at $500/month.

Pros: Viral growth potential. Large funnel. Try before you buy.

Cons: Free users cost money (infrastructure + support). Typical free-to-paid conversion is 2-4%. Attracts tire-kickers.

Verdict: Better for horizontal products. B2B buyers will pay for value upfront.

Usage-Based Pricing

Structure: $0.50 per customer scored/month. $5 per recommendation generated. No base fee.

Pros: Perfect alignment (pay for what you use). Low entry barrier. Scales with their growth.

Cons: Unpredictable revenue. Creates weird incentives (we want to score more, they want to score less). Complicated pricing.

Verdict: Could work at scale but too unpredictable for early stage.

Professional Services

Structure: $50-100K consulting engagements. Custom builds. $10K/month managed service.

Pros: High ACV. Deep relationships. Learn the problem intimately. Profitable immediately.

Cons: Doesn't scale (people-heavy). Long sales cycles. Different team needed. Hard to transition to product later.

Verdict: Good for first 5-10 customers to learn. Not sustainable.

Platform + Marketplace

Structure: Core platform at $500-1K/month. Marketplace for third-party playbooks/apps (20-30% take rate).

Pros: Two revenue streams. Ecosystem creates moat. Network effects.

Cons: Chicken-and-egg. Need customers to attract consultants. Complex two-sided marketplace. Requires major investment.

Verdict: Interesting for Year 3-5. Not viable for MVP.

Recommended Approach

Start with what we've designed: Base subscription + performance bonuses. Proven model, predictable revenue, shows confidence, scales well.

Year 2-3: Add professional services for Tier 3 customers (custom integrations, model tuning, consulting). 5-10% of revenue.

Year 3-4: Add marketplace (third-party playbook library, integration ecosystem, certified consultants). 10-15% of revenue.

Year 4+: Consider usage-based option for very large customers (10K+ records). Enterprise alternative.

Don't overcomplicate early. Nail the core model first.

Risks & Mitigation

Timing too early? Maybe growth-stage SaaS isn't ready for this level of sophistication.

- Likelihood: Low. CS platforms like Gainsight, ChurnZero, and Vitally all include churn-prediction features, indicating market demand. Need to figure out whether companies that can't afford these solutions feel the pain.
- Mitigation: Target companies are already trying to solve this (spreadsheets → us)

Incumbents move downmarket? Gainsight launches "Gainsight Lite" at $500/month.

- Likelihood: Medium (12-18 months out)
- Mitigation: 100+ customers before they move. Better self-service UX. Enterprise companies' track record moving downmarket is typically not good. Different business model.

Market consolidation? Salesforce acquires Gainsight, bundles for free.

- Likelihood: Low (already happened, didn't kill market)
- Mitigation: We're specialized. Better churn prediction.

Technical Risks

ML accuracy disappointing? Models under 70% accurate, not useful.

- Likelihood: Low (proven 85%+ in literature)
- Mitigation: Start with proven algorithms. Benchmark vs customer spreadsheets.

Data quality issues? Customer data too messy, missing key signals.

- Likelihood: Medium-High (real problem at SMBs)
- Mitigation: Data quality scoring at signup. Reject poor-quality trials. Educate customers.

Integration maintenance burden? APIs change, connectors break, customers churn.

- Likelihood: Medium
- Mitigation: Use Airbyte (community maintains). Monitor changes. SLA for critical integrations.

GTM Risks

CAC higher than projected? Actually $8-10K instead of $5K.

- Likelihood: Medium
- Mitigation: Focus organic/inbound first. Delay paid ads. Optimize conversion.

Sales cycle longer? 60-90 days instead of 30-45.

- Likelihood: Medium
- Mitigation: Improve trial experience. Reduce friction. Month-to-month contracts.

Churn higher than expected? 70-80% retention instead of 90%.

- Likelihood: Medium (SMB has higher churn)
- Mitigation: Customer success focus. “Meta” problem (we prevent churn).

Product Risks

“Sleeping dog” false positives? We flag low-usage customers as at-risk, reach out, and remind them they’re paying for something they don’t use. They cancel. We caused the churn.

- Likelihood: High (if we’re not careful)
- Mitigation: Train models on behavioral changes, not absolute usage. Distinguish “dormant

& stable" from "dormant & declining." Exclude stable dormant accounts from alerts and performance bonuses.

Customers don't act on recommendations? Great predictions, but CS teams ignore playbooks.

- Likelihood: Medium
- Mitigation: Track usage. Gamify ("You saved 12 customers!"). Integrate into workflows.

"Black box" problem? Customers don't trust ML predictions.

- Likelihood: Medium
- Mitigation: Explainability layer (show why). Let them override.

Performance bonus disputes? Disagreements about who we "saved."

- Likelihood: Medium-High
- Mitigation: Crystal-clear contracts. Transparent reporting. Audit rights. Build trust.

Execution Risks

Can't hire ML engineers? Competitive talent market.

- Likelihood: High
- Mitigation: Remote-first (global talent). Competitive comp. Meaningful equity.

Burn rate exceeds plan? Slower growth, run out of cash.

- Likelihood: High (most startups)
- Mitigation: Raise enough upfront (24-month runway). Monthly reviews. Cut burn if needed.

Founder burnout? 3+ years to profitability is a long grind.

- Likelihood: Medium
- Mitigation: Co-founder partnership. Advisory support. Work-life balance.

Bottom line: Lots of risks. But none are insurmountable. Execution matters more than perfect planning.

Conclusion

Why This Works

The problem is real and urgent. 35% churn is killing growth-stage SaaS companies. Every CEO knows it. Every board meeting discusses it.

The ROI is obvious. Even 5% churn reduction = 200%+ return. Easy to justify.

The market is underserved. Enterprise platforms are too expensive and too complex. Growth-stage companies need something built for them.

The technology exists. Airbyte, scikit-learn, Metabase - all proven, open source, ready to use. We're not inventing anything. We're assembling the pieces.

The business model aligns incentives. Performance bonuses mean we only win when customers win. Risk-sharing.

It scales. Software + ML scales infinitely. Not constrained by people.

Timing is right. Customer success is now a mature function. Budget exists. Buyers are educated. Just need a product that fits.

Next 90 Days

Weeks 1-4: Validate

- Interview 20 VPs of Customer Success
- Confirm pain, pricing, features
- Build financial model
- Create pitch deck

Weeks 5-8: Proof of Concept

- Set up Airbyte + basic model
- Demo to 5 prospects
- Get feedback

Weeks 9-12: MVP

- Hire founding engineer
- Build core pipeline + prediction engine + dashboard
- Sign 3-5 alpha customers (free for 6 months)

Months 4-6: Alpha

- Onboard alphas
- Iterate
- Achieve 85%+ accuracy
- Document 2-3 saves
- Raise $700K-1M from angels

Then we're off to the races.

Contact: [Your Email]

Date: November 2025

Let's build this. s: - Pre-written with personalization tokens - "Hi {{FirstName}}, I noticed you haven't tried {{FeatureName}} yet…" - Customer just clicks "Send"

Call discussion guides: - 3-5 talking points per at-risk customer - "Ask about: their experience with X, whether they're achieving Y goal, roadblocks to adoption" - Objection handling: common reasons for churn + responses

Intervention timeline: - "Day 1: Send email A" - "Day 3: If no response, call" - "Day 7: Follow-up email B" - "Day 14: Schedule QBR"

Customer exclusion: - Let customers remove accounts from churn list (maybe they have inside info we don't) - Only provide recommendations for remaining accounts - Track exclusion

reasons for our learning

Performance tracking: - Which customers were flagged high-risk? - Which ones survived 4+ months? - Calculate performance bonuses owed - Quarterly reconciliation report

Phase 3: Advanced ML + strategic insights (Months 6-12)

Better models: - Add Random Forest + XGBoost to ensemble - Weighted average of predictions - Target 90%+ accuracy - A/B test model versions

Quarterly strategic reports: - Cohort analysis: “Customers from Channel X churn at 2x rate vs Channel Y” - Product gaps: “Top 3 missing features cited by churned customers” - Training needs: “Help doc Y has 3x lower engagement than doc Z” - Support patterns: “Ticket type A predicts 60% churn probability”

These insights help fix systemic issues, not just save individual customers.

More integrations: - Expand to 15 connectors (Pendo, Amplitude, Helpscout, Drift, etc.) - CSV upload for custom data - REST API for programmable access

The Technical Stack

Data pipeline: - **Airbyte** (open source): 600+ connectors, handles schema drift, CDC support - **dbt**: SQL transformations, feature engineering - **PostgreSQL** or **Snowflake/BigQuery**: Data warehouse (depends on customer size) - Per-tenant schema isolation for security

ML layer: - **Python** with scikit-learn, XGBoost, LightGBM - **MLflow**: Experiment tracking, model versioning - Batch scoring: Daily runs, results stored in warehouse - REST API for on-demand predictions

Dashboard: - **Metabase** (open source): Embedded, white-labeled at higher tiers - Customer-facing UI: React frontend - Backend: Python (FastAPI)

Orchestration: - **Apache Airflow** or **Prefect**: Schedule jobs, coordinate pipelines

Hosting: - AWS or GCP, Kubernetes for auto-scaling - Separate S3 buckets per tenant for data sovereignty

Multi-Tenant Architecture

We’re not giving every customer their own infrastructure. That doesn’t scale.

Shared infrastructure, logical separation: - All customers on same Kubernetes cluster -

Shared Airbyte instance, shared ML compute - Separate schemas in the data warehouse (tenant*1.*, tenant*2.*) - Row-level security where needed - Separate S3 buckets for sensitive data

Why this works: - Cost-efficient (shared compute) - Fast onboarding (provision new schema in seconds) - Easy upgrades (apply to all tenants simultaneously) - Can move to database-per-tenant for enterprise customers if needed

Cost Per Customer

Infrastructure costs: - Data warehouse: $50-150/month (depends on data volume) - Airbyte Cloud (if not self-hosted): $100-200/month - Compute (ML training/scoring): $20-50/month - Metabase Cloud (if not self-hosted): $0-20/month - **Total: $170-420/month per customer**

At $800-2,000/month subscription, that's 75-80% gross margin. Good SaaS economics.

Open Source Advantage

Why we're using open source everywhere we can:

Lower costs: Airbyte Community Edition is free. PostgreSQL is free. scikit-learn is free. We're paying for hosting, not software licenses.

No vendor lock-in: Customer doesn't like us? They can export their data warehouse and build on it themselves. We're not holding anything hostage.

Battle-tested: These tools run at thousands of companies. They're proven, stable, well-documented.

Community support: Something breaks? Google it, find 10 Stack Overflow answers. vs proprietary tools where you're at the vendor's mercy.

Extensibility: Customer needs a custom connector? We can build it. vs closed-source tools where you have to beg the vendor for features.

The trade-off: We have to maintain infrastructure ourselves. But that's fine - we're a tech company. We should be good at ops.

Business Model & Pricing

The Tiers

Tier 1: Starter - $200/month

For small CS teams just getting started: - Up to 500 customer records - 3 data integrations (Stripe + CRM + one more) - Daily churn scoring - Email alerts - 50 recommendations/month ($2 per additional) - Email support (48hr response) - Performance bonus eligible

Target customer: Early growth-stage, 1-2 person CS team, limited budget

Tier 2: Professional - $800/month

This is where most customers will land: - Up to 2,000 customer records - 8 data integrations - Advanced ML (ensemble models) - Slack integration - Quarterly strategic reports - 200 recommendations/month ($1.50 per additional) - Priority support (email + chat, 24hr response) - Performance bonus eligible

Target customer: Core growth-stage companies, 3-6 person CS team

Tier 3: Enterprise - $2,000/month

For larger customers or those who want white-glove: - Up to 10,000 records - Unlimited integrations - White-labeled dashboards (remove our branding) - 1 custom integration per quarter - Monthly strategic reports - Unlimited recommendations - Dedicated success manager - Phone support (8hr response) - Quarterly business review - Performance bonus eligible

Target customer: Upper end of growth-stage, 7+ person CS team, approaching mid-market

Tier 4: Custom

Call us. Includes: - Database-per-tenant isolation (instead of shared) - SSO/SAML - Custom SLAs - On-premise deployment - API access - Whatever else you need

Performance Bonus Structure

Here's how it works:

Trigger: We flag a customer with ≥70% churn risk score.

Success: That customer is still active 4 months later.

Payment: We get 10% of that customer's ACV, one time.

Frequency limit: Max once per customer per 12 months (prevents double-counting).

Customer control: They can remove accounts from our prediction list (maybe they have inside info). Those exclusions aren't eligible for bonuses.

Example: - Company has 1,000 customers, $12K average ACV - We flag 30 as high-risk in Q1 - 20 of those survive 4+ months - Performance bonus: 20 × $12K × 10% = $24K - Customer saved: 20 × $12K × 70% margin = $168K gross profit - Our total fees: $24K subscription + $24K bonus = $48K - Their net benefit: $120K (250% ROI)

Why 4 months? - Too short (1-2 months) and we're not really proving we saved them - Too long (6+ months) and we're waiting forever to get paid - 4 months feels right - enough time to prove impact, not so long that customers forget

How do we avoid disputes? Clear contract language. Pre-agreed measurement methodology. Quarterly reconciliation report showing: - Customer ID - Flagged date - Risk score - ACV at time of flagging - Status at 4-month mark - Bonus due (if applicable)

Customer has audit rights. They can verify the list. If there's a disagreement, we discuss it. Relationship matters more than squeezing every dollar.

Add-On Revenue

Extra integrations: $100/month per connector beyond tier limits (only for standard Airbyte connectors)

Custom connector development: $2,000 upfront + $200/month maintenance (for specialized tools we don't have connectors for)

Professional services: - Data source consultation: $200/hour (helping them clean up their data) - Custom model tuning: $5,000 project (training models on their specific use case) - Strategic CS consulting: $300/hour (Tier 1-2 only; included in Tier 3)

API access: - Read API (export predictions): $500/month - Write API (push custom data): $1,000/month

We're not building a big professional services business. These are just options for customers who want extra help. Maybe 5-10% of revenue.

Unit Economics

Customer Acquisition Cost: $5,000

Breakdown: - Marketing spend: $1,500 (content, ads, demand gen) - Sales time: $2,000 (demos, trial support, closing) - Trial costs: $500 (infrastructure during trial period) - Onboarding: $1,000 (activation support)

That's blended average. Organic/inbound customers might be $3K. Outbound might be $7K. We're estimating $5K overall.

Lifetime Value:

Average customer at $1,200/month (mix of Tier 2-3): - Monthly revenue: $1,200 - Gross margin: 75% - Gross profit/month: $900 - Expected lifetime: 33 months (90% annual retention = 10% monthly churn) - LTV: $900 × 33 = $29,700

LTV:CAC = 5.9:1

That's healthy. SaaS benchmark is 3:1 minimum. We're at 6:1.

Payback period: 5.6 months

$5,000 CAC / $900 monthly gross profit = 5.6 months

Benchmark is <12 months good, <18 months acceptable. We're at 5.6 months. Excellent.

Why does this work? - Low infrastructure costs (open source, shared infrastructure) - Self-service reduces sales overhead - High retention creates long LTV - Performance bonuses are mostly margin (minimal cost to us)

Go-To-Market Strategy

Sales Model: Sales-Assisted PLG

We're not pure product-led growth. Deal size ($2.4K-24K/year) justifies some sales touch. Complex product (data + ML) benefits from human help. VP Customer Success is making a bet-your-job decision - they want to talk to someone.

But we're also not doing traditional enterprise sales. No custom demos, no 3-month evaluation cycles, no RFP processes. Customers can complete the entire journey online. Pricing is fixed (no negotiation for Tiers 1-3).

The hybrid approach:

Most leads (70%) start self-service: - Read our content, sign up for free trial - Automated onboarding emails - In-app guidance for connecting data sources - Dashboard activates after first data sync - “Book time with expert” button if they get stuck

But we touch every trial: - Day 3: “How’s it going? Any blockers?” (15-min call) - Day 7: “Let me show you advanced features” (20-min demo) - Day 14: “Ready to move forward?” (30-min close call)

Top 20% of leads get white-glove treatment: - Alpha/beta customers - Reference account candidates - Bigger deal size potential ($24K+) - We’ll do extra onboarding support, custom walkthroughs, QBRs

Acquisition Channels

Year 1: Content + Inbound

We’re not spending big on paid ads yet. Build organic engine first.

SEO content targeting CS leaders: - “How to Calculate Customer Churn Rate” - “10 Early Warning Signs Your Customer Is About to Leave” - “Build vs Buy: Customer Success Software Decision Framework” - Target: 10,000 organic visits/month by Month 12

Founder-led content: - LinkedIn posts (3x/week) on churn reduction tactics - Guest posts on CS blogs (ChurnKey, Custify, Gainsight’s blog) - Speaking at conferences (Customer Success Summit, Pulse)

Free tools as lead magnets: - Churn rate calculator - CAC payback calculator - Customer health score template spreadsheet - Gated downloads generate emails

Budget: $3K/month for content writers, SEO tools Target: 50 qualified leads/month by Month 12

Year 2: Add Paid Acquisition

Once organic is working, layer on paid:

Google Ads: - “Churn prediction software,” “customer success platform,” “reduce SaaS churn” - Budget: $5K/month - Target: 20 leads/month at $250 CPA

LinkedIn Ads: - Target job titles: VP Customer Success, Director CS, CCO - Target companies: $5-50M ARR B2B SaaS - Budget: $3K/month - Target: 10 leads/month at $300 CPA

Retargeting: - Website visitors who viewed pricing - Blog readers (3+ articles) - Budget: $2K/month

Total paid: $10K/month for 30 additional qualified leads/month

Year 2-3: Partnerships

CS consultants: - Find 10-20 independent consultants who advise growth-stage SaaS companies - Offer 15-20% recurring commission on referrals - They handle complex implementations, we provide the platform - Target: 20% of customers via partners by Year 3

Integration partnerships: - Co-marketing with Slack, Intercom, Zendesk, Segment - Get listed in their app marketplaces - Target: 10% of customers via integrations by Year 3

Community sponsorships: - Sponsor Customer Success Slack groups - Free accounts for CS influencers - Co-create case studies

The Sales Process

Typical journey (45 days):

Week 1: Awareness - Prospect finds us via blog post or Google search - Downloads “Churn Reduction Playbook” (gated) - Enters email nurture sequence

Weeks 2-4: Consideration - Gets 5 educational emails over 3 weeks - Reads comparison articles (us vs Gainsight vs ChurnZero vs build in-house) - Talks to their VP CS peers (word-of-mouth)

Weeks 5-6: Trial (14 days) - Signs up for free trial - Day 1-2: Connects data sources (self-service) - Day 3: Sales call #1 - “Any blockers? Let me help” - Days 3-7: Historical data loads, models train - Day 7: First predictions appear - Day 7: Sales call #2 - “Here’s what we’re seeing” - Days 8-14: Customer tests recommendations, shows to team - Day 14: Sales call #3 - “Ready to go forward?”

Week 7-8: Close - Tier 1-2: Online checkout (credit card) - Tier 3: Annual invoice (NET 30 terms) - Month-to-month for <$24K/year - Annual contract (10% discount) for $24K+

Weeks 8-11: Onboarding - Tier 1-2: Email onboarding, help docs - Tier 3: Kick-off call with dedicated success manager - Goal: First recommendation acted on within 30 days

Conversion metrics we’re targeting:

- Trial signup → Paid: 20% (industry benchmark is 15-25%)
- Sales cycle: 30-45 days
- Average deal size: $1,200/month ($14,400/year)

- Close rate of qualified leads: 30%

Marketing Budget

Year 1: $73K - Content: $36K - Tools: $12K - Events: $10K - Website: $15K

Year 2: $241K - Content: $48K - Paid ads: $120K - Tools: $18K - Events: $25K - Partnerships: $30K

Year 3: $324K - Content: $60K - Paid ads: $180K - Tools: $24K - Events: $40K - Partnerships: $20K

That gives us blended CAC of $5K (marketing + sales costs together).

Reality Check: Customer Business Failures

The Blindspot Nobody Wants to Talk About

The industry talks about 35% annual churn like it's all fixable. It's not.

Here's what the data shows: **92% of SaaS startups fail within 3 years**. 45% fail by year 5. 65% fail by year 10. Even after 10 years in business, only 13% of SaaS startups reach $10M ARR.

Our target customers are growth-stage companies ($5-50M ARR). They've survived early-stage, so they're not in the "92% fail in 3 years" bucket. But they're still vulnerable. 67% of SaaS companies experience their first major scaling crisis between $5M-$15M ARR.

Realistic Churn Model

Industry baseline (without us): 35% annual churn broken down as: - **Customer business failures: 8%** - Companies going out of business, getting acquired, pivoting away - **Controllable churn: 27%** - Product dissatisfaction, switching to competitors, budget cuts, poor onboarding

With our platform: 18% annual churn broken down as: - **Customer business failures: 8%** - Still happens, outside our control - **Controllable churn: 10%** - **We reduce this by 63%**

That's the value prop: **We can't save customers whose businesses fail, but we're experts at saving everyone else.**

Updated Unit Economics

With 18% annual churn (82% retention): - **Expected customer lifetime: 5.6 years** - **LTV: $46,000** - **CAC: $5,000** - **LTV:CAC ratio: 9.2:1** (excellent) - **CAC payback: 7.2 months**

This is healthy SaaS economics. The 9:1 LTV:CAC ratio gives us room for mistakes and shows the business model works even with realistic churn.

Financial Projections (Revised)

The Moderate Growth Path

With 18% annual churn, here's what it takes to reach ~200 customers by Year 3:

Year	New Customers	Churned	Ending Customers	Ending ARR
1	40	0	40	$442K
2	100	7	133	$1,468K
3	100	23	209	$2,307K

What this means: - We acquire 240 total customers over 3 years to end with 209 - We lose 31 customers to churn (mostly in Y2-Y3 as the base grows) - Total CAC spend: $1.2M ($5K × 240 customers) - Blended MRR: $920/month

Year 1

Revenue: - Average customers through year: 20 (ramping from 0 to 40) - Subscription revenue: $221K - Performance bonuses: $33K - Total revenue: $254K

Expenses: - COGS (25%): $64K - **Gross profit: $190K (75% margin)** - Sales & Marketing: $223K - Engineering: $540K (3 FTEs @ $180K) - Product: $160K (1 FTE) - CS/Support: $100K (1 FTE) - G&A: $32K - **Operating expenses: $1,055K** - **EBITDA: -$865K**

Burn rate: $72K/month

Year 2

Revenue: - Average customers through year: 87 (from 40 to 133) - Subscription revenue: $954K - Performance bonuses: $143K - Total revenue: $1,098K

Expenses: - COGS (25%): $274K - **Gross profit: $824K (75% margin)** - Sales & Marketing: $591K - Engineering: $900K (5 FTEs) - Product: $320K (2 FTEs) - CS/Support: $200K (2 FTEs) - G&A: $138K - **Operating expenses: $2,149K** - **EBITDA: -$1,325K**

Burn rate: $110K/month
Customers lost to churn: 7 (18% of 40)

Year 3

Revenue: - Average customers through year: 171 (from 133 to 209) - Subscription revenue: $1,888K - Performance bonuses: $283K - Total revenue: $2,171K

Expenses: - COGS (25%): $543K - **Gross profit: $1,628K (75% margin)** - Sales & Marketing: $824K - Engineering: $1,440K (8 FTEs) - Product: $480K (3 FTEs) - CS/Support: $400K (4 FTEs) - G&A: $314K - **Operating expenses: $3,458K** - **EBITDA: -$1,830K**

Burn rate: $152K/month
Customers lost to churn: 23 (18% of 133)

3-Year Summary

Metric	Year 1	Year 2	Year 3
Customers	40	133	209
ARR	$442K	$1,468K	$2,307K
Revenue	$254K	$1,098K	$2,171K
Gross Profit	$190K	$824K	$1,628K
EBITDA	-$865K	-$1,325K	-$1,830K
Cumulative Burn	$865K	$2,190K	$4,020K

Key metrics: - Total customers acquired: 240 - Total CAC spend: $1.2M - Cumulative 3-year burn: $4.0M - Net Revenue Retention: ~100% (82% retention + some upsells)

Path to Profitability

With 18% annual churn, breakeven happens in **Q4 Year 5**:

Year 4 projections: - Add 150 new customers, lose 38 to churn (18% of 209) - Ending: 321

customers - ARR: $3.5M - EBITDA: -$2.3M (still burning)

Year 5 projections: - Add 150 new customers, lose 58 to churn (18% of 321) - Ending: 413 customers - ARR: $4.6M - Revenue: $5.3M
- Gross profit: $4.0M - Operating expenses: $4.0M - **EBITDA: breakeven**

From Year 6 onwards, we're profitable and can grow without raising more capital.

Total funding needed: ~$7M - Seed (Year 1): $1.5M - Series A (Year 2-3): $3.0M - Series B (Year 4-5): $2.5M

This is realistic for a B2B SaaS company with strong unit economics (9:1 LTV:CAC). The 5-year path to profitability is longer than we'd like, but it reflects the reality of customer business failures we can't control.

Alternative Business Models We Considered

Pure Performance-Based (No Base Fee)

Structure: No subscription. We only charge 20-25% of retained revenue from saved customers.

Pros: Completely de-risks purchase for customer. Perfect incentive alignment. Premium positioning.

Cons: Zero revenue for first 4-6 months while proving value. Complex attribution. Hard to forecast. Needs $5M+ funding for 18-month runway.

Verdict: Too risky for bootstrap. Maybe viable later with more capital.

Freemium

Structure: Free tier for 100 customers, basic features. Paid tiers start at $500/month.

Pros: Viral growth potential. Large funnel. Try before you buy.

Cons: Free users cost money (infrastructure + support). Typical free-to-paid conversion is 2-4%. Attracts tire-kickers.

Verdict: Better for horizontal products. B2B buyers will pay for value upfront.

Usage-Based Pricing

Structure: $0.50 per customer scored/month. $5 per recommendation generated. No base fee.

Pros: Perfect alignment (pay for what you use). Low entry barrier. Scales with their growth.

Cons: Unpredictable revenue. Creates weird incentives (we want to score more, they want to score less). Complicated pricing.

Verdict: Could work at scale but too unpredictable for early stage.

Professional Services

Structure: $50-100K consulting engagements. Custom builds. $10K/month managed service.

Pros: High ACV. Deep relationships. Learn the problem intimately. Profitable immediately.

Cons: Doesn't scale (people-heavy). Long sales cycles. Different team needed. Hard to transition to product later.

Verdict: Good for first 5-10 customers to learn. Not sustainable.

Platform + Marketplace

Structure: Core platform at $500-1K/month. Marketplace for third-party playbooks/apps (20-30% take rate).

Pros: Two revenue streams. Ecosystem creates moat. Network effects.

Cons: Chicken-and-egg. Need customers to attract consultants. Complex two-sided marketplace. Requires major investment.

Verdict: Interesting for Year 3-5. Not viable for MVP.

Recommended Approach

Start with what we've designed: Base subscription + performance bonuses. Proven model, predictable revenue, shows confidence, scales well.

Year 2-3: Add professional services for Tier 3 customers (custom integrations, model tuning, consulting). 5-10% of revenue.

Year 3-4: Add marketplace (third-party playbook library, integration ecosystem, certified consultants). 10-15% of revenue.

Year 4+: Consider usage-based option for very large customers (10K+ records). Enterprise alternative.

Don't overcomplicate early. Nail the core model first.

Risks & Mitigation

Market Risks

Timing too early? Maybe growth-stage SaaS isn't ready for this level of sophistication. - Likelihood: Low (46% already using churn prediction) - Mitigation: Target companies already trying to solve this (spreadsheets → us)

Incumbents move downmarket? Gainsight launches "Gainsight Lite" at $500/month. - Likelihood: Medium (12-18 months out) - Mitigation: 100+ customers before they move. Better self-service UX.

Market consolidation? Salesforce acquires Gainsight, bundles for free. - Likelihood: Low (already happened, didn't kill market) - Mitigation: We're specialized. Better churn prediction.

Technical Risks

ML accuracy disappointing? Models under 70% accurate, not useful. - Likelihood: Low (proven 85%+ in literature) - Mitigation: Start with proven algorithms. Benchmark vs customer spreadsheets.

Data quality issues? Customer data too messy, missing key signals. - Likelihood: Medium-High (real problem at SMBs) - Mitigation: Data quality scoring at signup. Reject poor-quality trials. Educate customers.

Integration maintenance burden? APIs change, connectors break, customers churn. - Likelihood: Medium - Mitigation: Use Airbyte (community maintains). Monitor changes. SLA for critical integrations.

GTM Risks

CAC higher than projected? Actually $8-10K instead of $5K. - Likelihood: Medium - Mitigation: Focus organic/inbound first. Delay paid ads. Optimize conversion.

Sales cycle longer? 60-90 days instead of 30-45. - Likelihood: Medium - Mitigation: Improve

trial experience. Reduce friction. Month-to-month contracts.

Churn higher than expected? 70-80% retention instead of 90%. - Likelihood: Medium (SMB has higher churn) - Mitigation: Customer success focus. “Meta” problem (we prevent churn).

Product Risks

“Sleeping dog” false positives? We flag low-usage customers as at-risk, reach out, and remind them they’re paying for something they don’t use. They cancel. We caused the churn. - Likelihood: High (if we’re not careful) - Mitigation: Train models on behavioral *changes*, not absolute usage. Distinguish “dormant & stable” from “dormant & declining.” Exclude stable dormant accounts from alerts and performance bonuses.

Customers don’t act on recommendations? Great predictions, but CS teams ignore playbooks. - Likelihood: Medium - Mitigation: Track usage. Gamify (“You saved 12 customers!”). Integrate into workflows.

“Black box” problem? Customers don’t trust ML predictions. - Likelihood: Medium - Mitigation: Explainability layer (show why). Let them override.

Performance bonus disputes? Disagreements about who we “saved.” - Likelihood: Medium-High - Mitigation: Crystal-clear contracts. Transparent reporting. Audit rights. Build trust.

Execution Risks

Can’t hire ML engineers? Competitive talent market. - Likelihood: High - Mitigation: Remote-first (global talent). Competitive comp. Meaningful equity.

Burn rate exceeds plan? Slower growth, run out of cash. - Likelihood: High (most startups) - Mitigation: Raise enough upfront (24-month runway). Monthly reviews. Cut burn if needed.

Founder burnout? 3+ years to profitability is a long grind. - Likelihood: Medium - Mitigation: Co-founder partnership. Advisory support. Work-life balance.

Bottom line: Lots of risks. But none are insurmountable. Execution matters more than perfect planning.

Conclusion

Why This Works

The problem is real and urgent. 35% churn is killing growth-stage SaaS companies. Every CEO knows it. Every board meeting discusses it.

The ROI is obvious. Even 5% churn reduction = 200%+ return. Easy to justify.

The market is underserved. Enterprise platforms are too expensive and too complex. Growth-stage companies need something built for them.

The technology exists. Airbyte, scikit-learn, Metabase - all proven, open source, ready to use. We're not inventing anything. We're assembling the pieces.

The business model aligns incentives. Performance bonuses mean we only win when customers win. Risk-sharing.

It scales. Software + ML scales infinitely. Not constrained by people.

Timing is right. Customer success is now a mature function. Budget exists. Buyers are educated. Just need a product that fits.

Next 90 Days

Weeks 1-4: Validate - Interview 20 VPs of Customer Success - Confirm pain, pricing, features - Build financial model - Create pitch deck

Weeks 5-8: Proof of Concept - Set up Airbyte + basic model - Demo to 5 prospects - Get feedback

Weeks 9-12: MVP - Hire founding engineer - Build core pipeline + prediction engine + dashboard - Sign 3-5 alpha customers (free for 6 months)

Months 4-6: Alpha - Onboard alphas - Iterate - Achieve 85%+ accuracy - Document 2-3 saves - Raise $700K-1M from angels

Then we're off to the races.

Contact: [Your Email]
Date: November 2025

Let's build this.