

Kyanos Interview

Ed Forman in Conversation with Claude of Anthropic
(November 6, 2025)

What got you thinking about AI and progressive politics?

I've gotten to know Scott Anderson of Strategic Victory Fund through friends in Sun Valley, where I spend about half of each year. Over breakfast this summer, Scott asked a question that's stuck with me: *"What threats to progressives are on the horizon that we're not seeing?"*

We'd been commiserating about how everyone missed the impact of podcasts. Then it hit me — the podcasts of 2028 are going to be chatbots. And left to their own devices, those chatbots are not likely to promote progressive candidates or causes.

Since that morning, I've been a little obsessed with understanding how AI systems shape public perception — and what we can do to make sure truth and democratic values don't get crowded out by convenience or bias.

You've spent most of your career in tech — how does that shape how you see this?

I spent most of my career in the trenches of Silicon Valley, managing and helping to build very early-stage startups. Over about thirty-five years, I was an early employee at eleven of them — each one a front-row seat to how technology gets created, hyped, and transformed. Not one of those companies still does business under the name it had when I joined it, which tells you something about how volatile those early waves can be.

I've always been drawn to the next big thing — sometimes too early. I've been there for decision-support systems, SQL databases, client-server computing, wireless email, desktop analytics, and interactive media. Many times, I was there before the market was ready.

So you've been living inside these shifts for decades. Does that make you more optimistic or worried about AI in politics?

A little of both. Technology always moves faster than institutions can adapt. My career has been

a reminder that whoever organizes information best — wins. That's what scares me about this moment.

How did you move from building startups to teaching and coaching?

About fifteen years ago, I felt the need for a reset. Stanford launched a program called SEED, pairing experienced businesspeople with small and medium-sized companies in developing economies. I signed up as a coach, and it was transformative. My wife — who's also a Stanford MBA — and I lived in Ghana for most of a year and coached companies in Ghana, Nigeria, and Côte d'Ivoire.

It was a life-changing experience and became the launching pad for my second career: teaching entrepreneurship and business-model design at Stanford for mid-career executives and advising interesting companies, mostly run by former students.

That period also changed how I think about systems — economic or political — and what happens when you don't build capacity early.

When you talk about chatbots becoming the new information gateways, what kind of impact do you actually foresee?

It's hard to overstate it. We're heading into a world where most voters won't "Google" a candidate — they'll ask an AI. And whatever answer comes back will feel factual, neutral, even authoritative. The challenge is that AI models don't judge ideas — they train on what's easiest for them to find, parse, and connect.

Recent research shows this shift is happening fast. Google's AI Overviews already appear on roughly half of all search results. **That shift took less than a year — a structural change in how people discover information.**

If progressive ideas aren't structured for this new reality, they risk being invisible when it matters most.

So this isn't about partisan bias — it's about how information gets organized?

Exactly. This isn't "SEO on steroids." It's a whole new set of rules. Search engines ranked pages; AI systems interpret relationships — between facts, entities, and claims. If the data about progressive ideas isn't structured for that, it doesn't just get buried — it effectively doesn't exist to the model.

The right has spent years building deep, linked networks of policy content, think tanks, and media that are easy for machines to ingest. Progressives have plenty of truth and evidence, but it's scattered, unstructured, and inconsistent. We've been publishing for people — not for machines that now mediate what people see.

It's not ideology that drives the imbalance — it's information architecture.

I've been reading about partisan chatbots — is that what you're proposing?

Not exactly. I just wrote a short white paper on that. There's a surge of partisan chatbots coming — conservatives already have several in development. Progressives probably need at least one of their own. The New York Times just ran a piece about that trend.

But that's not what Kyanos is. We're the factual layer underneath all that — the source of structured, evidence-based information that any chatbot can draw from. Even the most partisan chatbot still needs a credible source.

Why is what Kyanos proposes to do hard?

First, because the SEO era only had to optimize for one surface: Google's PageRank. The AI era has many — ChatGPT, Gemini, Perplexity, Copilot — each with its own algorithms, formats, and update cycles.

Then there's the workflow problem. Campaign operatives and communications professionals are used to writing for persuasion, not for machines.

That's not intuitive. It requires retraining habits that have been built over decades of political communication — habits optimized for persuasion, not precision. Getting seasoned operatives to write and publish with this dual audience in mind — people and machines — is going to be a significant behavioral shift.

And building tools that support the creation of machine-readable content is going to be a significant user-experience design challenge. How do the writer and editor work with AI in the content development process? That's a behavioral shift —

and it won't happen without tools that make it natural.

How do you reach scale quickly?

Our go-to-market strategy leans heavily on partnerships — roughly 75 percent of distribution will come through platform integrations. The two dominant campaign-website systems, WordPress and NationBuilder, together power about half of Democratic campaign sites.

By integrating with them, we make structured content a default, not an extra chore. Partner distribution is key: they bring trust and customers; we bring the technology and support.

We'll also run a lightweight certification program for agencies and consultants — think of it as "*structured-data-ready*" training. Once they're certified, their work naturally feeds accurate information into the AI ecosystem.

What does Kyanos actually sell, and to whom?

Campaigns, officeholders, committees, and causes — anyone who depends on being seen accurately in AI-driven discovery.

The customers pay for clarity. They subscribe to a SaaS platform that audits how their content appears in generative systems, suggests markup fixes, and validates performance.

Between election cycles, the same tools serve officeholders — monitoring how policies, quotes, and records appear in AI answers and flagging distortions before they spread. **That continuous "officeholder mode" keeps the lights on between campaigns and turns visibility itself into infrastructure.**

How do you price this, and what do the economics look like?

We'll operate on a subscription model — annual plans scaled to race size or organization scope, with optional managed-service tiers.

The economics are strong because the heavy lift is in building and maintaining the schema and the monitoring system. Once the data engine runs, each new customer adds incremental revenue with limited servicing cost.

That structure also makes it easier to prove ROI — both politically and financially. When a

campaign can literally *see* that its issues and bio are surfacing accurately across AI assistants, that's measurable value.

What's the market potential here?

The Democratic infrastructure market — including campaigns, PACs, advocacy groups, and aligned organizations — totals roughly \$600–800 million a cycle. Within that, digital communication and data optimization make up about \$150 million.

Our initial Serviceable Available Market (SAM) is around \$90 million — the slice of that infrastructure likely to adopt AI-discovery tools between 2025 and 2028. Our Serviceable Obtainable Market (SOM) — realistic early customers — is about \$6 million by 2028, or roughly 259 clients on subscription.

That's a small share of the total spend, but strategically critical — because whoever builds the “discovery layer” first sets the standard for everyone else. **This is a network-effect business disguised as an infrastructure tool.**

What are the biggest risks — and how do you mitigate them?

Three big categories.

First, **platform risk**: we depend on access to multiple AI systems that evolve constantly. The fix is diversification — never relying on a single model or API, and documenting every query method so that results remain reproducible.

Second, **legal and compliance risk**: campaign-finance rules are strict about coordination and disclosure. We'll keep tight legal separation between any nonprofit partnerships and our commercial work, advised by counsel steeped in election law.

Third, **adoption risk**: the behavioral change. That's why we're working through CMS partners and agencies people already trust — to meet campaigns where they are.

Transparency is our insurance policy — every piece of structured content carries a provenance trail that shows where it came from and when it was verified.

Why structure Kyanos as a for-profit B-Corp instead of a nonprofit?

Because sustainability matters. Nonprofits depend on grants; campaigns need reliability.

A B-Corp lets us take in investment capital, reward efficiency, and still be mission-locked. It aligns social impact and capital discipline: if we deliver trustworthy visibility for progressives, everyone wins.

We'll also keep an affiliated nonprofit arm for civic research and transparency work, but it won't touch campaign activity. The separation protects both sides legally and ethically. **The B-Corp gives us flexibility to act fast while being held accountable to purpose, not partisanship.**

Who needs to be on the founding team?

Three core leaders.

The **CEO** sets direction and alignment — someone who can speak to investors and campaign professionals with equal fluency, translating urgency into action.

The **CTO** builds the trust layer — ensuring our tools are credible, secure, and interoperable with the broader progressive tech stack.

The **Chief Customer Officer** drives adoption and outcomes — connecting the product to field operations, PAC partners, and communication agencies.

What matters most is that they trust each other enough to move fast — and know when to argue and when to ship.

And how will governance work?

We'll maintain a small, balanced board — impact investors, movement leaders, and an independent technologist to keep product integrity front and center.

The structure enforces dual accountability: to mission and to capital. The CEO is the only officer with a board seat, ensuring clarity of command without compromising transparency.

That mix keeps us from tilting too far toward ideology or pure venture logic. Decisions stay grounded in visibility, credibility, and democratic resilience.

When investor interests and progressive needs don't line up, how do you make the call?

It's built into the charter. As a B-Corp, we're legally required to balance profit and purpose. If a decision strengthens visibility, credibility, and trust in progressive information ecosystems, it's good business.

If it doesn't, it isn't.

That clarity saves time. We're not trying to maximize exits — we're trying to maximize democratic resilience.

Politics is cyclical. What happens between election seasons?

That's where our "officeholder mode" matters.

After elections, the same platform helps elected officials maintain an accurate public record in AI environments — tracking how their statements, votes, and policies are represented.

It turns AI visibility from a campaign asset into ongoing democratic infrastructure.

That recurring use gives us a stable subscription base and constant data feedback to refine the models before the next cycle.

Why can't this wait until next July for the next accelerator cohort?

Because the calendar doesn't wait for us.

Tech takes time to design, test, and deploy.

If we want validated results from the 2026 cycle — the proof we'll need for broad acceptance in 2028 — the work must start now.

AI systems learn continuously. The 2028 narrative is already being written by today's training data.

So what does success actually look like?

Success in 2026 means measurable evidence that campaigns and causes using Kyanos are consistently visible, accurate, and equitably represented in AI answers.

That's our proof of concept.

By 2028, success means this becomes infrastructure — a shared layer of the progressive data ecosystem, with hundreds of campaigns and officeholders using a common schema and the SCOPE monitoring system keeping watch for bias across major AI platforms.

Proof in 2026. Scale in 2028. Build it once, prove it works, then make it part of how democracy protects itself in the AI age.

And what kind of investment does it take to get there?

In the near term, the critical capital is the kind that lets us assemble the core team — enough runway to attract and retain the right technical and product people and give them confidence this is a funded, focused effort.

Once that team is in place, we refine the plan, validate assumptions, and confirm the cost structure.

A business plan is a set of hypotheses — pricing, customer acquisition, data-monitoring cost — all of which need validation before scaling.

This isn't a long runway; 2026 is already around the corner.

We're building the plane while it's rolling down the runway.

We can't afford a waterfall development cycle — the market and technology move too fast.

We need a cockpit crew and we need to get started.

And your role in all this?

My startup days are behind me. I don't want to run another company.

What I do want is to help get this initiative built — to assemble the right people, frame the problem clearly, and design a structure that can outlast any one individual.

I'd be thrilled to leave this conversation knowing that, through HGL's investments, there are already well-resourced teams tackling these problems.

The role I see for myself is as an advisor — helping put the company together, identifying key team members, getting them up to speed quickly, being a cheerleader and a sounding board.

I won't take an operational role, but I want to make sure democracy has a seat inside the new information infrastructure before it's written without us.